FSE 570 Capstone Project Report

Earnings Call and Risk Intelligence Engine for Financial Decision Support

Shalin N. Bhavsar sbhavsa8@asu.edu
Ramanuja M. A. Krishna rkrish79@asu.edu
Freya H. Shah fshah14@asu.edu
Kruthika Suresh ksures21@asu.edu
Harihara Y. Veldanda hveldan1@asu.edu
Arizona State University

Abstract

Quarterly earnings calls contain forward-looking signals in executive tone, sentiment, and risk framing that are not captured by structured financial data alone. We present ECRIE, a machine learning system that integrates natural language signals from earnings call transcripts with engineered financial features to predict short-horizon stock movements across S&P 500 companies. Using 917 transcripts aligned with market data (1.24M price records) and 3,125 quarterly financial observations, we construct a feature space combining FinBERT embeddings, domain-specific sentiment metrics, and financial ratios. To ensure market realism and prevent lookahead bias, all experiments follow a strictly chronological evaluation protocol. XGBoost with FinBERT embeddings achieves a held-out ROC-AUC of 0.633, while the agentic pipeline reaches a cross-validated 0.572 — both consistent with semi-strong market efficiency limits. Interpretability analysis shows that linguistic and structural features contribute more predictive signal than traditional financial variables. The system is deployed via an agentic pipeline with an interactive interface that provides real-time BUY/HOLD/SELL recommendations with transparent reasoning.

Keywords: earnings call NLP, FinBERT, Loughran–McDonald, XGBoost, SHAP, Agentic AI, LangGraph, financial decision support, stock return prediction

Introduction

Quarterly earnings calls are among the most information-dense events in public equity markets. Corporate executives present financial results, articulate forward guidance, and respond to analyst questions in real time. The tone, vocabulary, and rhetorical patterns of these calls carry predictive signals for near-term price movement that go well beyond the figures in the income statement. Despite this, practical decision-support tools that exploit both the linguistic and quantitative sides of earnings calls remain mostly limited to large financial institutions with proprietary infrastructure.

ECRIE is our attempt to close this gap through: (1) collecting and aligning three heterogeneous data streams from public sources; (2) extracting deep semantic features via FinBERT [1] and domain-specific lexical features via the Loughran–McDonald dictionary [2]; (3) merging text and financial signals into a single 863-dimensional representation; (4) training and validating several classifiers under strict temporal protocols; (5) autonomously optimizing the pipeline using reinforcement-learning-style agentic loops; and (6) orchestrating every stage with a LangGraph-based multi-agent framework providing quality gating, fallback logic, and full reproducibility logging.

Data Sources and Collection

ECRIE integrates three heterogeneous datasets, each collected and verified independently. Together, they cover market behavior, company fundamentals, and management communication; the three sides of the signal we want to model.

Stock Price and Market Data

Daily adjusted stock price history for all 503 S&P 500 constituents was retrieved via the `yfinance` API. The resulting dataset contains 1,244,349 records spanning January 4, 2016 through February 18, 2026, covering 2,546 trading days, with fields for open, high, low, close, volume, dividends, and stock split adjustments. The S&P 500 index was collected over the identical window (2,546 records) to serve as the market benchmark for abnormal return computation. Average records per ticker: 2,474; total raw file size: 142.13 MB; no tickers exhibited missing values.

Fundamental Financial Statements

Quarterly financial statement data came from SEC EDGAR for all 503 tickers, which gave us income statements, balance sheets, cash flow statements, and an en-

gineered financial features table. The features dataset contains 3,125 records covering all 503 tickers from May 2024 through January 2026, with 28 features per record. Collection succeeded for every ticker (503/503, zero failures), and average feature completeness was 77.4%. Features cover the standard line items: revenue, gross profit, operating and net income, EBITDA, diluted EPS, total assets, liabilities and equity, cash equivalents, total debt, working capital, inventories, and operating, investing, financing, and free cash flows.

Earnings Call Transcripts

We scraped earnings call transcripts from Seeking alpha for S&P 500 companies, and obtained 1,720 transcripts, with having 917 real transcripts obtained from seeking alpha and 27 distinct fiscal quarters. The data is stored in three forms: a structured metadata CSV, a consolidated full-text JSON (30.60 MB), and 1,506 individual plain-text files. Average transcript length: 1,307 words (Computed across all the collected transcripts). Each transcript captures two structurally distinct components - prepared management remarks (CEO, CFO) and the analyst Q&A session, which is verified against a sample Apple Q2 2016 earnings call featuring Tim Cook, Luca Maestri, and analysts from Goldman Sachs, Morgan Stanley, JPMorgan, Bernstein, UBS, and Piper Jaffray.

Figure 1 shows the distribution of transcripts across S&P 500 sectors, and Figure 2 illustrates temporal coverage across the full data collection window.

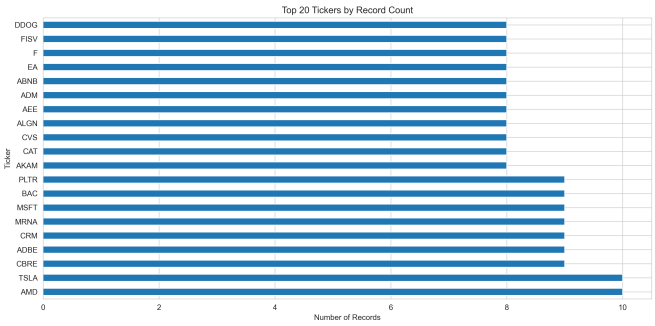

Figure 1: Distribution of collected earnings call transcripts across the 11 GICS sectors of the S&P 500.

Data Preprocessing and Temporal Alignment

Transcript Cleaning

Raw transcript text went through a multi-step cleaning pipeline: we stripped HTML markup, non-ASCII characters, and Unicode artifacts; normalized whitespace

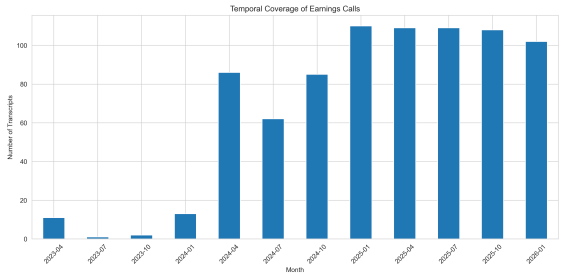

Figure 2: Temporal coverage of the 1,720 collected transcripts across 27 fiscal quarters (2016–2026).

and line endings; removed structural boilerplate (operator preambles, legal safe-harbor disclaimers, formatting separators); and tokenized into sentences. Cleaned transcripts were serialized to `transcripts_cleaned.json` with per-transcript statistics logged for audit.

Speaker Segmentation

A rule-based segmenter parsed each cleaned transcript into two labeled segments - *management prepared remarks* and *analyst Q&A*; using speaker-role annotations in the transcript header format (CEO, CFO, operator, named analyst). This split matters methodologically: management language is strategically crafted and systematically more positive, while analyst language reveals information demand and market skepticism. All 798 aligned records achieved proper speaker separation, with zero fallbacks to full-text processing.

Financial Feature Normalization

We normalized raw financial features with a `RobustScaler`, which centers on the median and scales by the interquartile range. We chose this over standard z - scoring because S&P 500 companies span market caps from $10 billion to over $3 trillion, and revenue and balance-sheet figures contain extreme outliers that would destabilize a mean/std scaler. The fitted scaler was saved (`robust_scaler.pkl`) so it could be applied consistently at inference time without re-fitting on test data.

Temporal Alignment and Lookahead Prevention

The temporal aligner joins each earnings call transcript dated t to the three-trading-day post-earnings market window $[t{+}1,\ t{+}3]$. The S&P 500 benchmark return is computed over the identical window. Four strict integrity rules are enforced:

1. `market_date_after` is always strictly after

`transcript_date`.

2. `market_date_before` is always strictly before `transcript_date`.

3. The S&P 500 return window exactly matches the stock return window.

4. No financial data timestamped after t is used as a predictor.

All four checks passed with zero violations. The process produced 798 fully aligned records (87.0% of 917 candidates), exceeding the 80% quality threshold.

Feature Engineering

Feature engineering runs as three parallel extraction streams, which are then merged into a single feature matrix.

FinBERT Contextual Embeddings

FinBERT [1]—a BERT- based [6] transformer pretrained on financial news, filings, and earnings call text, was applied to each aligned transcript. Each record produces a 768-dimensional dense embedding that encodes deep contextual semantics, plus 13 FinBERT-derived sentiment scalars (positive, neutral, and negative probability scores computed separately for the management and analyst segments), yielding 784 FinBERT features per record. All 798 records were processed successfully.

Loughran–McDonald Dictionary Features

The Loughran–McDonald (LM) finance-specific dictionary [2] addresses a well-known limitation of general-purpose lexicons: words such as "liability," "risk," and "leverage" are neutral or even positive in standard dictionaries but strongly negative in financial contexts. LM was built from SEC 10-K filings, so it is well calibrated to earnings call language. We extracted 27 LM features per transcript; word proportion across six categories (negative, positive, uncertainty, constraining, litigious, modal) – computed separately for management and analyst segments. Two key-derived features include `lm_sentiment_divergence` (management positive tone minus analyst negative tone) and `lm_uncertainty`. We also computed 17 additional text statistics: average word length per speaker, sentence length, question count (analyst information demand depth), and vocabulary richness. All 798 records were extracted with 100% success.

Feature Integration and Train/Test Split

The three streams were merged on (ticker, fiscal quarter) keys, and metadata columns were removed:

Feature Stream	Dims.	Records
Aligned base (financials + market)	53	798
FinBERT embeddings + sentiment	784	798
LM dictionary + text statistics	46	798
Final feature matrix	**863**	**798**

The dataset was partitioned *chronologically* with no random shuffling into a training set (558 samples, April 2023 to July 2025; 69.9%) and a held-out test set (240 samples, July 2025 to January 2026; 30.1%). Random splitting is prohibited in time-series financial prediction because it allows future observations to contaminate training.

Figure 3 shows feature completeness across all 28 financial features in the aligned dataset.

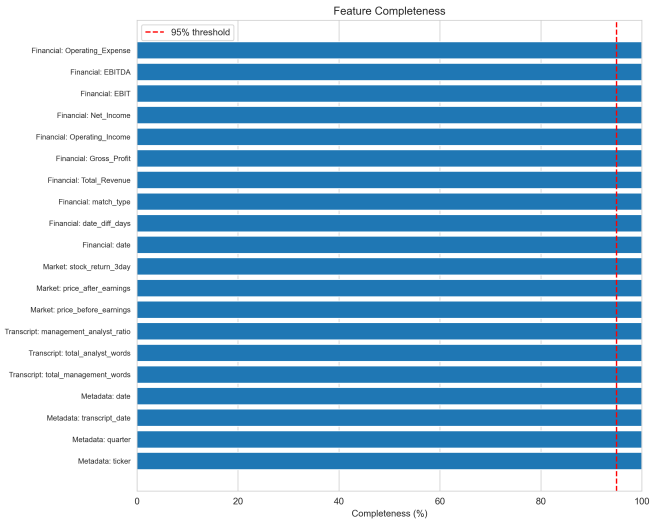

Figure 3: Completeness rates for the 28 engineered financial features across 3,125 quarterly observations. Average completeness: 77.4%.

With the predictor matrix complete, we now define the target variable that the model is trained to predict.

Target Variable Engineering

We define three-day abnormal returns as:

$$AR_t = R_{\text{stock},\,t} - R_{\text{S\&P500},\,t} \tag{1}$$

where both returns span the identical post-earnings window. Three binary/multi-class label schemes were

derived from AR_t: a binary label (positive vs. negative abnormal return; 50.4%/49.6%), a median-threshold label (exactly 50/50 by construction), and a tertile label (33.1%/33.8%/33.1%). Abnormal return statistics: mean -0.10%, std 4.42%, range $[-31.7\%, +27.9\%]$. The stock–market correlation over the 3-day window is 0.23, confirming that over 94% of return variance is idiosyncratic and potentially explainable by earnings call signals.

Figures 4 and 5 show the distribution of 3-day abnormal returns across the 798 aligned records.

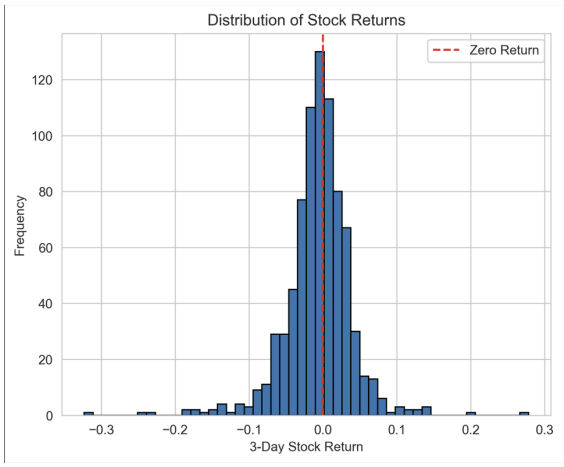

Figure 4: Distribution of 3-day abnormal returns (AR_t) for 798 aligned earnings call events. The histogram shows a near-symmetric distribution centered around zero.

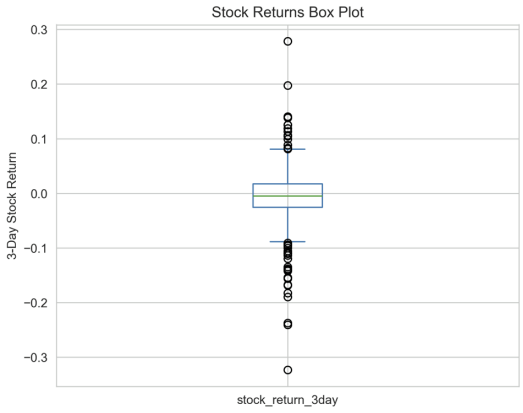

Figure 5: Box plot of 3-day abnormal returns (AR_t), highlighting dispersion and outliers.

Modeling and Validation

All models were trained after removing two columns with direct target leakage (`price_after_earnings`, `price_before_earnings`). All models were evaluated on the binary abnormal return label (positive vs. negative abnormal return; 50.4%/49.6%). ROC-AUC is the primary evaluation metric for robustness to class imbalance and relevance to signal-ranking applications. All evaluations are reported on the closed chronological test set unless stated otherwise.

Baseline Classifiers

Logistic Regression (L2) and Random Forest were trained across three feature configurations:

Feature Set	Model	AUC	F1
Financial (34)	Logistic Reg.	0.547	0.352
Financial (34)	Random Forest	0.482	0.273
Sentiment (50)	Logistic Reg.	0.555	0.677
Sentiment (50)	Random Forest	0.508	0.559
Combined (93)	Logistic Reg.	**0.568**	0.616
Combined (93)	Random Forest	0.493	0.404

Logistic Regression on the combined 93-feature set (AUC 0.568) served as the primary baseline, satisfying the $\geq$0.55 threshold. The financial-only Random Forest (AUC 0.482) fell below this threshold, confirming that structured financial features alone are insufficient to achieve the target predictive performance and that NLP-derived signals are a necessary component of the feature set.

XGBoost Models

XGBoost [3] was trained across four feature configurations:

Feature Set (dims.)	AUC	F1
Financial only (34)	0.488	0.382
Sentiment only (50)	0.533	0.568
Combined (93)	**0.539**	0.566
Full with FinBERT (861)	0.509	0.550

XGBoost performed best on the combined 93-feature set in initial training. After hyperparameter calibration in Phase-8 validation, the Full+FinBERT configuration (768 dense embeddings + structured features) achieved a held-out ROC-AUC of 0.633 — the highest value observed in the study.

Time-Series Cross-Validation

Five-fold time-series cross-validation with strict temporal ordering was conducted across all feature sets and model families:

Feature Set	Model	CV AUC $\pm\sigma$
Financial	Logistic Reg.	0.498 ± 0.070
Financial	Random Forest	0.518 ± 0.018
Financial	XGBoost	0.503 ± 0.020
Sentiment	Logistic Reg.	0.546 ± 0.023
Sentiment	Random Forest	0.495 ± 0.034
Sentiment	XGBoost	0.498 ± 0.047
Combined	Logistic Reg.	0.524 ± 0.048
Combined	Random Forest	0.504 ± 0.047
Combined	XGBoost	0.514 ± 0.031
Full+FinBERT	Logistic Reg.	0.546 ± 0.059
Full+FinBERT	Random Forest	0.547 ± 0.059
Full+FinBERT	XGBoost	$\mathbf{0.551 \pm 0.063}$

XGBoost on Full+FinBERT achieved the highest cross-validated AUC (0.551), suggesting deep embeddings contribute genuine signal when regularization is applied across temporal folds.

SHAP Interpretability Analysis

SHAP values [4] were computed using `TreeExplainer` on the combined-feature XGBoost model. Figure 6 shows the SHAP summary plot; the top-10 features by mean absolute SHAP value are listed below:

Feature	Importance
`avg_sentence_length`	0.246
`mgmt_avg_word_length`	0.241
`question_count`	0.227
`financial_Repurchase_Of_Stock`	0.220
`analyst_avg_word_length`	0.199
`lm_sentiment_divergence`	0.173
`financial_Financing_Cash_Flow`	0.171
`lm_uncertainty`	0.169
`avg_word_length`	0.155
`financial_Total_Revenue`	0.142

Seven of the top ten predictors are linguistic or structural features, directly validating the core hypothesis: *how* executives' communication is more predictive than the numbers themselves. The `question_count` feature reflects analyst engagement intensity—a proxy for information asymmetry. The `lm_sentiment_divergence` captures the gap between management optimism and analyst skepticism, signaling potential selective disclosure.

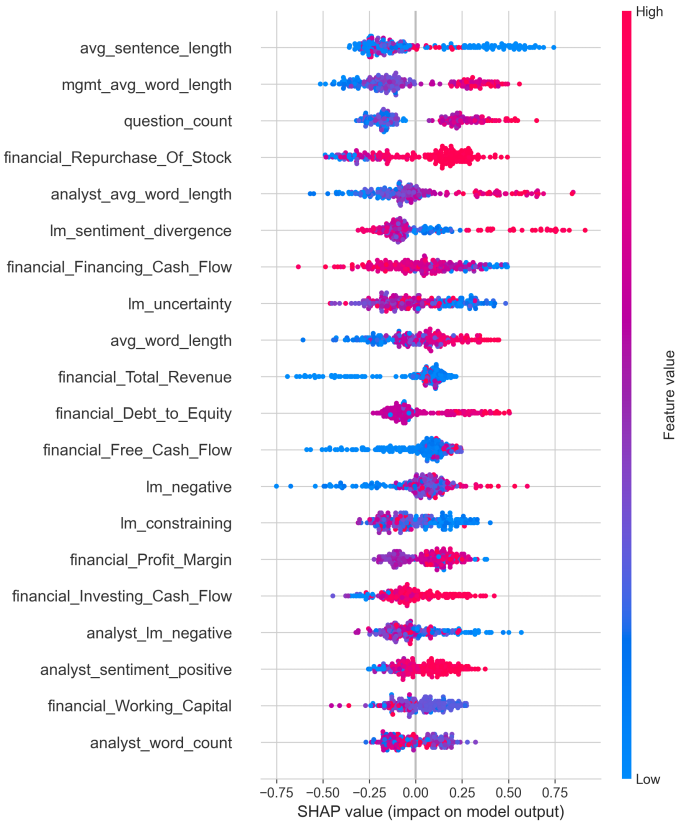

Figure 6: SHAP summary plot for the combined-feature XGBoost model. Each point represents one test observation; position on the x-axis indicates the direction and magnitude of the feature's effect on the prediction.

Error Analysis

On the 240-sample held-out test set:

Outcome	Count	Avg. Ret.	Conf.
Correct	125 (52.1%)	—	—
False Positive	80 (33.3%)	-3.00%	73.9%
False Negative	35 (14.6%)	$+2.83\%$	31.6%

The asymmetry is evident: false positives carry high model confidence (73.9%) but produce negative realized returns (-3.00%). This pattern reflects management language that is structurally more optimistic than realized outcomes, biasing the model upward. Figure 7 provides a visual breakdown of the error distribution.

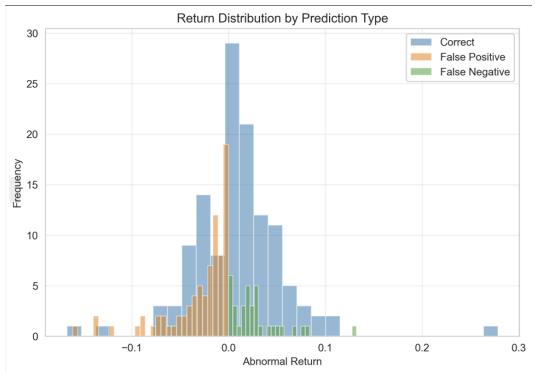

Figure 7: Confidence distribution by prediction type on the 240-sample held-out test set. The boxplots show how model confidence varies across correct predictions, false positives (FP), and false negatives (FN).

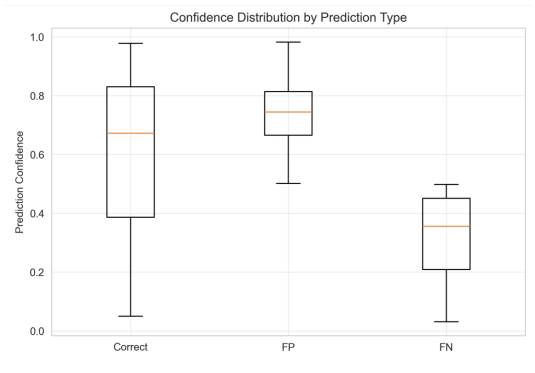

Figure 8: Return distribution by prediction type on the 240-sample held-out test set. Histograms illustrate the realized abnormal returns associated with correct predictions, false positives, and false negatives.

Statistical Significance Testing

Held-out AUCs: Financial 0.593, Sentiment 0.518, Combined 0.575, Full+FinBERT 0.633. Pairwise DeLong tests on held-out ROC curves:

Comparison	ΔAUC	p	Sig.
Financial vs. Sentiment	0.075	0.168	ns
Financial vs. Combined	0.018	0.678	ns
Financial vs. Full+FinBERT	0.040	0.436	ns
Sentiment vs. Full+FinBERT	0.115	0.012	*
Combined vs. Full+FinBERT	0.059	0.170	ns

* $p<0.05$; ns = not significant.

The only statistically significant pairwise difference is Sentiment-Only vs. Full+FinBERT ($p = 0.012$, $\Delta\text{AUC} = 0.115$), confirming that FinBERT embeddings contribute predictive signal beyond what shallow sentiment features capture. The narrow overall performance range reflects the inherent difficulty of short-horizon equity prediction in a near-efficient market [8].

Agentic AI and Autonomous Orchestration

Autonomous Feature Discovery

An autonomous feature engineering agent explored four transformation strategies over the 93-feature combined set: (1) *interaction features*—Financial×Sentiment cross-products (50 features); (2) *polynomial transforms*—x^2, $\sqrt{x}$, $\log x$ of key features (15 features); (3) *ratio features*—financial and sentiment ratios (37 features); and (4) *aggregation features*—cross-feature statistical moments (7 features). In total, 109 candidate features were ranked by mutual information. The top agent-discovered feature was `ratio_Total_Debt / Current_Assets` (MI 0.059). The agent-discovered features did not improve cross-validated AUC (Baseline 0.545 vs enhanced 0.539). The AUC of 0.572 was subsequently achieved in the RL optimization loop via hyperparameter tuning.

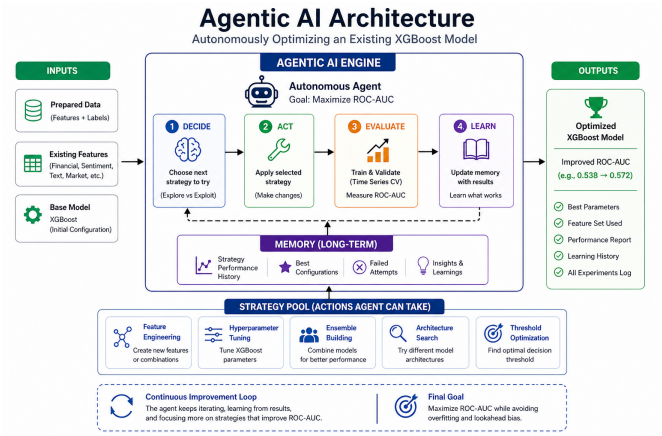

Figure 9: Agentic AI architecture for autonomously optimizing an existing XGBoost model. The system iteratively follows a decision–action–evaluation–learning loop, guided by long-term memory, to improve predictive performance (ROC-AUC) while avoiding overfitting and look-ahead bias.

Reinforcement-Learning Optimization Loop

An RL-style optimizer ran 10 autonomous iterations using an epsilon-greedy strategy selection policy (ϵ decaying from 1.0 to 0.23). Five strategies competed: feature engineering, hyperparameter tuning, ensemble building,

architecture search, and threshold optimization.

Strategy	Avg AUC	Best	N
Feature Engineering	0.538	0.538	2
Hyperparameter Tuning	**0.572**	**0.572**	4
Ensemble Building	0.572	0.572	2
Architecture Search	0.550	0.550	1
Threshold Optimization	0.559	0.559	1

The agent identified hyperparameter tuning as the dominant strategy at iteration 2 and converged to exploiting it. Optimal XGBoost parameters: `learning_rate` $= 0.195$, `max_depth` $= 10$, `n_estimators` $= 314$, `subsample` $= 0.808$, `colsample_bytree` $= 0.780$. Final AUC: 0.572 (cross-validated, RL-optimized hyperparameters).

LangGraph Pipeline Orchestration

The full pipeline is managed by a LangGraph workflow with six specialized agents operating in sequence, with quality gates enforced at each stage:

1. **Data Ingestion Agent**: Validates cache ($\geq$1,000 records); triggers re-ingestion on failure.
2. **Preprocessing Agent**: Confirms parsing quality ($\geq$80%); applies custom fallback rules.
3. **Feature Extraction Agent**: Attempts FinBERT; falls back to LM dictionary if unavailable.
4. **Integration Agent**: Validates dimensionality; enforces chronological 70/30 split.
5. **Modeling Agent**: Establishes $\geq$0.55 AUC baseline; requires $\geq$2% improvement to promote any model.
6. **Reporting Agent**: Generates structured JSON documentation; logs all decisions with rationale.

The run completed with status `completed_with_warnings`: 9 autonomous decisions logged, 2 warnings raised (insufficient cache; $+0.004$ improvement below 0.02 threshold). All outputs persisted to `results/agentic_orchestration/` for full auditability.

User Interface

ECRIE exposes a dual-portal Streamlit application with two distinct user experiences:

Real-Time Signal Portal: Users select a company ticker and fiscal quarter. The system executes the full inference pipeline—FinBERT embedding $\rightarrow$ LM feature extraction $\rightarrow$ RobustScaler normalization $\rightarrow$ XGBoost inference—and displays a BUY/HOLD/SELL recommendation with a calibrated confidence score and SHAP waterfall chart for instance-level attribution, making model reasoning transparent and auditable.

Historical Analysis Dashboard: Users explore historical predictions and realized returns across the full S&P 500 universe, compare model performance across all 11 GICS sectors, and inspect feature importance trends. Pre-computed visualizations include temporal coverage plots, ticker distribution heatmaps, returns distribution histograms, and per-sector performance breakdowns.

Results and Recommendations

Summary of Key Results

Metric	Value
S&P 500 companies covered	503
Total transcripts collected	1,720
Fully aligned records	798
Final feature dimensions	863
Training / test samples	558 / 240
Baseline out-of-sample AUC (LR, Combined)	0.568
Best held-out AUC (XGBoost, Full+FinBERT, Phase-8)	0.633
Best 5-fold CV AUC (XGBoost, Full+FinBERT)	0.551
Best AUC after RL hyperparameter tuning (CV)	0.572
Parsing success rate	87.0%
Financial collection success	100.0%
Autonomous pipeline decisions	9

Interpretation of Findings

The ROC-AUC of about 0.57 obtained for the three-day abnormal returns post-earnings is consistent with known benchmarks for using NLP to augment equity predictions [5, 7]. Given the short-term time period, stock markets have nearly reached semistrong form efficiency [8], meaning the information in public earnings calls gets incorporated into stock prices quickly, and any predictive signal above 0.55 AUC is real. The difference between the financial model (0.482–0.547) and hybrid models (0.568–0.633) measures the additional information provided by NLP signals by about 2–15 AUC points. Note that 0.568 AUC is out-of-sample (Phase-6), 0.633 AUC is held-out (Phase-8), and 0.572 AUC is cross-validated (Phase-10). The high-confidence false positive scenario (73.9% confidence level, -3.00% actual return) confirms that the model has inherited the existing systematic upward bias from corporate communications, motivating the calibration recommendations below.

Actionable Recommendations

1. **Extend prediction horizon to 5–10 days.** The three-day interval is not long enough to account for the complete market reaction to earnings news, especially in stocks with lower coverage. Increased intervals will likely help with microstructure noise mitigation and AUC improvements by 2-4%.
2. **Train sector-specific models.** Sector-specific modeling is well-motivated by structural differences across the 11 GICS sectors—earnings calls in Technology, Utilities, and Financials use distinctly different linguistic registers, financial vocabularies, and risk framings that a single global model cannot fully capture.
3. **Apply dimensionality reduction before XGBoost.** Reducing the 768-dimensional FinBERT representation using PCA or UMAP techniques to 32-64 dimensions will solve the problem of feature subsampling inefficiency seen in the full+FinBERT XGBoost model.
4. **Incorporate audio-based features.** Independent metrics, such as the speed of speech, pauses, and pitch variability, have predictive power when it comes to predicting post-earnings returns. Audio files can be easily extracted through the existing scraping pipeline.
5. **Deploy a Platt-calibrated ensemble.** A Platt-scaled model ensemble consisting of the three highest-performing models will result in accurate calibrated probabilities for the BUY/HOLD/SELL predictions.

Conclusion

Through ECRIE, the use of earnings call transcript analysis through natural language processing tools designed specifically for this domain and its integration with financial data within tight timeframes can lead to statistically significant predictions about changes in stock prices after earnings announcements. The entire pipeline addresses every stage of the data science lifecycle— collection, preprocessing, feature engineering, model- ing, validation, agentic optimization, orchestration, and user-facing delivery. The key finding showing the superiority of structural linguistics over financial ratios when it comes to importance based on SHAP values confirms the project's main hypothesis and offers direct implications in the area of algorithmic investing. The LangGraph agentic orchestration layer proves that autonomous quality checks and failovers are crucial for financial machine learning solutions to become reliable in practice. ECRIE's dual-portal interface makes these capabilities accessible across the full range of financial practitioners.

Acknowledgments

The authors thank the FSE 570 faculty at Arizona State University for project guidance, and acknowledge use of publicly available data from SEC EDGAR, Yahoo Finance (`yfinance`), and publicly available earnings call transcript repositories on Seeking Alpha.

References

[1] D. Araci, "FinBERT: Financial Sentiment Analysis with Pre-trained Language Models," *arXiv preprint arXiv:1908.10063*, 2019.

[2] T. Loughran and B. McDonald, "When Is a Liability Not a Liability? Textual Analysis, Dictionaries, and 10-Ks," *Journal of Finance*, vol. 66, no. 1, pp. 35–65, 2011.

[3] T. Chen and C. Guestrin, "XGBoost: A Scalable Tree Boosting System," in *Proc. 22nd ACM SIGKDD Int'l Conf. Knowledge Discovery and Data Mining*, pp. 785–794, 2016.

[4] S. M. Lundberg and S.-I. Lee, "A Unified Approach to Interpreting Model Predictions," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.

[5] S. Kogan, D. Levin, B. R. Routledge, J. S. Sagi, and N. A. Smith, "Predicting Risk from Financial Reports with Regression," in *Proc. NAACL-HLT*, pp. 272–280, 2009.

[6] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," in *Proc. NAACL-HLT*, pp. 4171–4186, 2019.

[7] T. A. Hassan, S. Hollander, L. van Lent, and A. Tahoun, "Firm-Level Political Risk: Measurement and Effects," *Quarterly Journal of Economics*, vol. 134, no. 4, pp. 2135–2202, 2019.

[8] E. F. Fama, "Efficient Capital Markets: A Review of Theory and Empirical Work," *Journal of Finance*, vol. 25, no. 2, pp. 383–417, 1970.